

six_classify

September 7, 2020

```
[1]: # Load the TensorBoard notebook extension
    %load_ext tensorboard

[2]: import tensorflow as tf
    from keras import layers
    from keras.preprocessing import image
    from keras.preprocessing.image import ImageDataGenerator
    import keras.backend as K
    K.set_image_data_format('channels_last')

    import numpy as np
    import matplotlib.pyplot as plt
    from matplotlib.pyplot import imshow
    import datetime
    from alexnet import AlexNet

[3]: # Set the GPU growth in order to avoid the sudden stop of the runtime.
    gpus = tf.config.experimental.list_physical_devices('GPU')
    for gpu in gpus:
        tf.config.experimental.set_memory_growth(gpu, True)

[4]: # Give the global constants. Please notify BATCH_SIZE for model.fit() and ↵
    ↵Batch_Size for
    # model.evaluate() and model.predict()
    EPOCHS = 50
    BATCH_SIZE = 32
    Batch_Size = 1
    image_width = 227
    image_height = 227
    channels = 3
    num_classes = 6

[5]: # Call the alexnet model in alexnet.py
    model = AlexNet((image_width, image_height, channels), num_classes)

[6]: # Compile the model
    model.compile(optimizer=tf.keras.optimizers.Adam(0.001),
```

```
loss='categorical_crossentropy',
metrics=['accuracy'])
```

```
[7]: # It will output the AlexNet model after executing the command
model.summary()
```

Model: "alex_net"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 55, 55, 96)	34944
max_pooling2d (MaxPooling2D)	(None, 27, 27, 96)	0
conv2d_1 (Conv2D)	(None, 27, 27, 256)	614656
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 256)	0
conv2d_2 (Conv2D)	(None, 13, 13, 384)	885120
conv2d_3 (Conv2D)	(None, 13, 13, 384)	1327488
conv2d_4 (Conv2D)	(None, 13, 13, 256)	884992
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 256)	0
flatten (Flatten)	(None, 9216)	0
dense (Dense)	(None, 4096)	37752832
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16781312
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 1000)	4097000
dense_3 (Dense)	(None, 6)	6006

Total params: 62,384,350
 Trainable params: 62,384,350
 Non-trainable params: 0

```
[8]: train_dir = '/home/mike/Documents/Six_Classify_AlexNet/seg_train/seg_train'
test_dir = '/home/mike/Documents/Six_Classify_AlexNet/seg_test/seg_test'
```

```
predict_dir = '/home/mike/Documents/Six_Classify_AlexNet/seg_pred/'
```

```
[9]: # keras.preprocessing.image.ImageDataGenerator
train_datagen = ImageDataGenerator(rescale=1.0/255)

# keras.preprocessing.image.DirectoryIterator
train_generator = train_datagen.flow_from_directory(train_dir,
                                                    ↵
                                                    ↪target_size=(image_width,image_height),
                                                    class_mode='categorical')

train_num = train_generator.samples
```

Found 14034 images belonging to 6 classes.

```
[10]: # Start Tensorboard --logdir logs/fit
log_dir="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
callback_list = [tensorboard_callback]
```

```
[11]: # Set verbose=1 (or verbose=0) for visibile (or invisible) training procedure.
model.fit(train_generator,
          epochs=EPOCHS,
          steps_per_epoch=train_num//BATCH_SIZE,
          callbacks=callback_list,
          verbose=1)
```

Epoch 1/50

438/438 [=====] - 24s 55ms/step - loss: 1.7000 - accuracy: 0.3986

Epoch 2/50

438/438 [=====] - 16s 36ms/step - loss: 1.0280 - accuracy: 0.6016

Epoch 3/50

438/438 [=====] - 16s 36ms/step - loss: 0.9226 - accuracy: 0.6450

Epoch 4/50

438/438 [=====] - 16s 36ms/step - loss: 0.8542 - accuracy: 0.6784

Epoch 5/50

438/438 [=====] - 16s 36ms/step - loss: 0.8028 - accuracy: 0.7041

Epoch 6/50

438/438 [=====] - 16s 36ms/step - loss: 0.7288 - accuracy: 0.7337

Epoch 7/50

438/438 [=====] - 16s 36ms/step - loss: 0.6962 -

```

accuracy: 0.7483
Epoch 8/50
438/438 [=====] - 16s 36ms/step - loss: 0.6328 -
accuracy: 0.7779
Epoch 9/50
438/438 [=====] - 16s 36ms/step - loss: 0.6266 -
accuracy: 0.7769
Epoch 10/50
438/438 [=====] - 16s 36ms/step - loss: 0.5795 -
accuracy: 0.7966
Epoch 11/50
438/438 [=====] - 16s 36ms/step - loss: 0.5597 -
accuracy: 0.8065
Epoch 12/50
438/438 [=====] - 16s 36ms/step - loss: 0.5405 -
accuracy: 0.8154
Epoch 13/50
438/438 [=====] - 16s 36ms/step - loss: 0.5115 -
accuracy: 0.8248
Epoch 14/50
438/438 [=====] - 16s 36ms/step - loss: 0.4910 -
accuracy: 0.8287
Epoch 15/50
438/438 [=====] - 16s 36ms/step - loss: 0.4928 -
accuracy: 0.8304
Epoch 16/50
438/438 [=====] - 16s 36ms/step - loss: 0.4638 -
accuracy: 0.8388
Epoch 17/50
438/438 [=====] - 16s 36ms/step - loss: 0.4366 -
accuracy: 0.8496
Epoch 18/50
438/438 [=====] - 16s 36ms/step - loss: 0.4207 -
accuracy: 0.8579
Epoch 19/50
438/438 [=====] - 16s 36ms/step - loss: 0.4173 -
accuracy: 0.8602
Epoch 20/50
438/438 [=====] - 16s 37ms/step - loss: 0.3991 -
accuracy: 0.8627
Epoch 21/50
438/438 [=====] - 16s 37ms/step - loss: 0.4020 -
accuracy: 0.8672
Epoch 22/50
438/438 [=====] - 16s 37ms/step - loss: 0.3782 -
accuracy: 0.8740
Epoch 23/50
438/438 [=====] - 16s 37ms/step - loss: 0.3824 -

```

accuracy: 0.8749
Epoch 24/50
438/438 [=====] - 16s 37ms/step - loss: 0.3312 -
accuracy: 0.8887
Epoch 25/50
438/438 [=====] - 17s 38ms/step - loss: 0.3652 -
accuracy: 0.8805
Epoch 26/50
438/438 [=====] - 16s 37ms/step - loss: 0.3286 -
accuracy: 0.8927
Epoch 27/50
438/438 [=====] - 16s 37ms/step - loss: 0.3036 -
accuracy: 0.9005
Epoch 28/50
438/438 [=====] - 16s 37ms/step - loss: 0.3273 -
accuracy: 0.8935
Epoch 29/50
438/438 [=====] - 17s 38ms/step - loss: 0.3107 -
accuracy: 0.8967
Epoch 30/50
438/438 [=====] - 16s 37ms/step - loss: 0.2917 -
accuracy: 0.9077
Epoch 31/50
438/438 [=====] - 17s 38ms/step - loss: 0.2435 -
accuracy: 0.9214
Epoch 32/50
438/438 [=====] - 17s 39ms/step - loss: 0.2651 -
accuracy: 0.9169
Epoch 33/50
438/438 [=====] - 17s 39ms/step - loss: 0.3002 -
accuracy: 0.9064
Epoch 34/50
438/438 [=====] - 17s 39ms/step - loss: 0.2608 -
accuracy: 0.9171
Epoch 35/50
438/438 [=====] - 17s 38ms/step - loss: 0.2831 -
accuracy: 0.9099
Epoch 36/50
438/438 [=====] - 17s 38ms/step - loss: 0.2871 -
accuracy: 0.9099
Epoch 37/50
438/438 [=====] - 16s 37ms/step - loss: 0.2332 -
accuracy: 0.9262
Epoch 38/50
438/438 [=====] - 16s 37ms/step - loss: 0.2171 -
accuracy: 0.9330
Epoch 39/50
438/438 [=====] - 16s 37ms/step - loss: 0.2182 -

```

accuracy: 0.9309
Epoch 40/50
438/438 [=====] - 16s 37ms/step - loss: 0.2403 -
accuracy: 0.9259
Epoch 41/50
438/438 [=====] - 17s 38ms/step - loss: 0.5131 -
accuracy: 0.8249
Epoch 42/50
438/438 [=====] - 17s 38ms/step - loss: 0.3985 -
accuracy: 0.8719
Epoch 43/50
438/438 [=====] - 16s 38ms/step - loss: 0.4105 -
accuracy: 0.8699
Epoch 44/50
438/438 [=====] - 16s 38ms/step - loss: 0.2782 -
accuracy: 0.9143
Epoch 45/50
438/438 [=====] - 17s 38ms/step - loss: 0.2710 -
accuracy: 0.9189
Epoch 46/50
438/438 [=====] - 17s 38ms/step - loss: 0.3167 -
accuracy: 0.9019
Epoch 47/50
438/438 [=====] - 17s 38ms/step - loss: 0.2383 -
accuracy: 0.9282
Epoch 48/50
438/438 [=====] - 16s 38ms/step - loss: 0.2565 -
accuracy: 0.9228
Epoch 49/50
438/438 [=====] - 16s 37ms/step - loss: 0.2467 -
accuracy: 0.9299
Epoch 50/50
438/438 [=====] - 16s 37ms/step - loss: 0.2256 -
accuracy: 0.9324

```

```
[11]: <tensorflow.python.keras.callbacks.History at 0x7f9f6024cad0>
```

```
[12]: %tensorboard --logdir logs/fit
```

```
<IPython.core.display.HTML object>
```

```
[13]: # It is the test generator as similar as the above.
test_datagen = ImageDataGenerator(rescale=1.0/255)

test_generator = test_datagen.flow_from_directory(test_dir,
                                                    ↵
                                                    ↪target_size=(image_width,image_height),
```

```
class_mode='categorical')
```

```
test_num = test_generator.samples
```

Found 3000 images belonging to 6 classes.

```
[14]: # Evaluate the trained model and return both the loss and the test accuracy.
evals = model.evaluate(test_generator,
                        verbose=1,
                        batch_size=Batch_Size,
                        steps=test_num//Batch_Size)

print("Loss = " + str(evals[0]))
print("Test Accuracy = " + str(evals[1]))
```

```
3000/3000 [=====] - 87s 29ms/step - loss: 0.9312 -
accuracy: 0.7851
Loss = 0.9311785697937012
Test Accuracy = 0.7850593328475952
```

```
[15]: # Give the implicit steps=7301 for selecting the specific image number.
predict_datagen = ImageDataGenerator(rescale=1.0/255)

predict_generator = predict_datagen.flow_from_directory(predict_dir,
                                                         ↪target_size=(image_width,image_height),
                                                         batch_size=Batch_Size,
                                                         ↪class_mode='categorical')


predict_num = predict_generator.samples
```

Found 7301 images belonging to 1 classes.

```
[16]: # Make the prediction for any one of the predicted images
predictions = model.predict(predict_generator,
                             verbose=1,
                             batch_size=Batch_Size,
                             steps=predict_num//Batch_Size)
```

```
7301/7301 [=====] - 23s 3ms/step
```

```
[17]: # Plot the descriptive diagram
imshow(predict_generator[256][0][0])
plt.imsave("predicted1.png",predict_generator[256][0][0])
```



```
[18]: predictions[256]
```


```
[18]: array([9.9994230e-01, 4.7316188e-15, 1.1871999e-12, 8.8631305e-14,  
        1.2097377e-13, 5.7670841e-05], dtype=float32)
```

```
[19]: print(predictions[256])
```

```
[9.9994230e-01 4.7316188e-15 1.1871999e-12 8.8631305e-14 1.2097377e-13  
5.7670841e-05]
```

```
[20]: imshow(predict_generator[1024][0][0])
```

```
[20]: <matplotlib.image.AxesImage at 0x7f9f100e5750>
```

```
[21]: predictions[1024]
```

```
[21]: array([3.5577407e-01, 7.1741693e-02, 3.4085230e-04, 1.7436147e-04,
          3.0319974e-05, 5.7193863e-01], dtype=float32)
```

```
[22]: import os

def get_category(predicted_output):
    return os.listdir(train_dir)[np.argmax(predicted_output)]
```

```
[23]: print(get_category(predictions[2048]))
```

forest

```
[24]: fig, axs = plt.subplots(2, 3, figsize=(10,10))

axs[0][0].imshow(predict_generator[32][0][0])
axs[0][0].set_title(get_category(predictions[32]))

axs[0][1].imshow(predict_generator[64][0][0])
axs[0][1].set_title(get_category(predictions[64]))

axs[0][2].imshow(predict_generator[128][0][0])
axs[0][2].set_title(get_category(predictions[128]))
```

```

axs[1][0].imshow(predict_generator[512][0][0])
axs[1][0].set_title(get_category(predictions[512]))

axs[1][1].imshow(predict_generator[1000][0][0])
axs[1][1].set_title(get_category(predictions[1000]))

axs[1][2].imshow(predict_generator[2000][0][0])
axs[1][2].set_title(get_category(predictions[2000]))

```

[24]: Text(0.5, 1.0, 'street')



[]: